Ethical Foundations and Responsibilities in eXplainable AI

Prof. Charlie Abela

Department of Artificial Intelligence
University of Malta

Explainable AI: Definition and Scope

Explainable AI (XAI) is a toolbox of methods that make model decisions understandable to people, but explainability lives on a spectrum rather than being all-or-nothing. What counts as "understandable" depends on audience, context, and purpose.

Always ask three framing questions: explainable to whom, for what purpose, and at what level of detail? These questions highlight that explanation quality is a socio-technical choice shaped by human factors and value judgements, not just algorithms.

Stakeholder Needs

Engineers want calibrated feature attributions, clinicians need clinically meaningful stories, and loan applicants require clear, fair reasons that enable contestation.

Cognitive Limits

Technically complete outputs can fail if they exceed human cognitive limits. A 100-feature attribution may be "right" but unusable, whilst a top-5 summary can be more actionable.

Design Trade-offs

Designing explanation systems inevitably privileges some needs over others. Optimising for trust may sacrifice debugging precision.

Intrinsic vs Post-Hoc Explainability

Intrinsic Interpretability

Intrinsically interpretable models are readable by design, such as linear regression, decision trees with if-then rules, and GAMs with visualisable shape functions. These give exact stories about their own mechanics with no approximation gap.

Post-Hoc Methods

Post-hoc explainability treats the trained model as a black box and explains after training using tools like SHAP and LIME. SHAP distributes predictions via Shapley values, whilst LIME fits a local surrogate to mimic the model near a point.

Fidelity Distinction

Intrinsic methods describe the model exactly, whereas post-hoc methods approximate and can misrepresent complex interactions. A sparse linear surrogate may explain itself, not the true non-linear boundary.

Industry Practice

Industry often favours post-hoc explanations for high-capacity models, assuming complexity is necessary for performance. This assumption should be tested empirically against strong interpretable baselines.

Trade-off Considerations

Choosing between approaches involves trade-offs in expressiveness, computational overhead, and trustworthiness. In high-stakes domains, prove black-box complexity is empirically necessary before sacrificing interpretability.

Revisiting the Accuracy-Interpretability Trade-Off

> The standard narrative claims you must choose: accuracy from opaque models or interpretability from simpler ones. That framing is often overstated and can mislead design choices.

Evidence shows well-designed interpretable models can rival black boxes in healthcare risk scoring, recidivism prediction, and fraud detection. GAMs and sparse rules have matched or approached ensembles and deep nets in multiple studies.

1 Engineering Convenience

The persistence of the "trade-off" can reflect engineering convenience rather than necessity. AutoML and transfer learning make black boxes easy to ship.

2 Ethical Question

We must ask whether black boxes are chosen for superior outcomes or developer convenience. In domains affecting rights and welfare, picking opacity when interpretable models suffice is an ethical failure.

3 Rigorous Evaluation

Compare baselines on held-out generalisation, subgroup robustness, distribution shift, and inference efficiency–not just i.i.d. accuracy. Invert the burden of proof: prove you need the black box.

Explainability's Purposes Beyond Trust

Explanations serve multiple critical functions in AI systems, extending far beyond simply building user confidence. Each purpose demands different approaches and levels of detail.

Verification

Explanations surface whether models learned meaningful signals rather than artefacts. Imaging models might latch onto scanner marks; credit models might use postcode as a proxy–issues visible only through explanation.

Debugging

They accelerate debugging by revealing failure modes, miscalibration, and counterintuitive interactions tied to domain knowledge. This enables targeted fixes such as new data collection or feature redesign.

Legal Compliance

Explanations are often a legal requirement for accountability and contestation, e.g., GDPR's demand for "meaningful information" and lending's adverse-action duties. Explainability is part of due process.

Domain Discovery

They can teach us about the domain by highlighting novel relationships or risk factors, turning models into hypothesis generators. Insights can seed new clinical studies or socioeconomic investigations.

Important consideration: These aims can conflict. Trust-oriented narratives may hide technical flaws, whilst engineer-grade details can overwhelm lay users. Effective practice tailors explanation type and granularity to the stakeholder and task.

Limits of Explanation and When to Pause

Some systems encode high-dimensional, non-linear, or stochastic structure that resists human-scale narratives. Summaries can clarify, but beyond a point they distort reality and invite overconfidence. Intellectual humility is essential.

Selective Narratives

Every explanation is selective and embeds value judgements about what to highlight. Even accurate selections can mislead by downplaying factors important to other stakeholders.

Explanation Gaming

Systems can satisfy formal "explanation" requirements whilst masking problematic logic. Post-hoc methods can rationalise arbitrary behaviour with plausible stories, enabling ethics theatre if unchecked.

Cognitive Constraints

Human cognition struggles with many interacting variables and conditional probabilities. Design for cognitive limits–progressive disclosure and top-k factors–whilst acknowledging residual uncertainty.

Prediction vs Explanation

Prediction and explanation can diverge when strong predictors are not causal levers, e.g., spurious correlates that boost accuracy. In high-stakes settings, limit or avoid deployment when faithful, usable explanations are not possible.